Comparative Analysis of Performance Using Server-Client Protocols

Mihail Costea, Liviu Chircu
University Politehnica of Bucharest
Faculty of Automatic Control and Computer Science
Emails: mihail.costea90@gmail.com, liviu.chircu@gmail.com

Abstract—Current web applications' solutions for bi-directional communication are widely based on AJAX. Even though they are well documented solutions that are backed up by years of utilization, they have limitations imposed by the HTTP protocol. HTTP is a stateless protocol that requires each connection to be treated as a new connection, requiring an unnecessary overhead to communicate in both directions. Because of these limitations, a new solution was developed - WebSockets - which are able to natively mantain a bi-directional channel, thus reducing the overhead imposed by connection establishment.

This paper proposes to exemplify the advantages and disadvantages between traditional HTTP implementations for bi-directional communication based on AJAX and WebSockets. It also presents the comparison of two WebSockets implementations, and proposes an architecture for a testing platform of different WebSockets implementations based on the obtained results.

I. INTRODUCTION

Along with the introduction of Web 2.0, web applications have been able to modify the content of HTML documents without the need of a complete page refresh, offering a more interactive environment for the end-user. The most popular technology used to create this interactivity is AJAX [1] (Asynchronous JavaScript and XML). With AJAX, pages can be updated in real-time, without needing an explicit action from the user. But because AJAX is based on HTTP - a stateless protocol that requires each connection to be treated as a new connection - bi-directional channels between 2 devices can only be simulated through a series of methods which impose a significant amount of overhead, as both nodes need to mimic this channel that require states. Bi-directional channels are necessary for web applications such as instant messaging clients, browser-based games, video calls, etc., where large amounts of traffic (with complete sets of headers and message bodies) need to be sent from each endpoint over to the opposite one.

The most popular solutions used by AJAX are: **polling**, where a client sends a request to a server at regular intervals with the server responding immediately and closing the connection; **long-polling**, where the client requests information from the server just as in normal polling, but the server now maintains its connection open until it has the data to send, point in time where it will send an immediate response; **streaming**, where the connection between the client and the server is kept alive indefinitely and data is streamed until one of the two closes the connection. The problem with last solution is that AJAX appends the new data to previously sent data until the connection ends, which is unnecessary in many cases. Moreover, HTTP headers must be sent whenever a new connection is created, which is a common case for polling and long-polling. If a great amount of small pieces of data are to be sent, as in the case of a web-based chat application, the overhead of the headers might actually overweigh the data that is actually transmitted [2].

In order to overcome the limitations of AJAX-based solutions, the HTTP-based WebSockets protocol was created [3]. HTTP was chosen as a base because most firewalls allow HTTP and HTTPS traffic (ports 80 and 443), while other ports' traffic might be blocked. Although based on HTTP, WebSockets do not inherit its limitations. HTTP is only used to create and close the connection between client and server, while the in-between information exchange is done through a series of WebSocket-specific data frames. The distinction between normal HTTP traffic and WebSockets is done through the "Upgrade: websocket" / "Connection: Upgrade" headers [3].

As mentioned above, WebSockets are a recently created technology (the RFC was written in 2011), so they are not supported by all major web servers, let alone being extensively tested. This may very well be considered a downside when compared to AJAX, which benefits from years of testing and industry usage. Although a series of testing platforms with the purpose of validating WebSockets implementations have been created, such as Autobahn [4], to our best knowledge, we have yet to find any performance-oriented testing platforms for a given WebSockets implementation.

The first part of paper compares AJAX with WebSockets. It is then followed by a comparison of two WebSockets implementations based on *Node.js*. Based on the comparison results, it proposes two different architectures for a testing platform capable of automating the testing process. One is independent of the underlying Operating System and present devices. It should work for both laptops and smartphones, for Windows and Linux and other devices and Operating Systems. The other one is platform dependent, but it could be used not only for testing how WebSockets behave in general, but also to

assess the behaviour of specific applications which are based on WebSockets.

II. RELATED WORK

Puranik et al. [5] exemplifies the differences in terms of performance between AJAX and WebSockets. The authors added real-time monitoring support for OASIS [6], an open-source real-time instrumentation middleware for distributed real-time and embedded systems. The collected instrumentation data was sent over the Web using AJAX and WebSockets and the results were compared. The WebSockets server consumes 50% less network bandwidth than the AJAX server. Also, the WebSockets client consumes memory at a constant rate, while the AJAX client consumes it at an increasing rate. Furthermore, the WebSockets-oriented implementation can send up to 215.44% more data samples than the AJAX-based counterpart although it demands the same amount of network bandwidth.

Even though WebSockets are better than AJAX based solutions, they are still not as good as raw TCP sockets. Agarwal [7] discusses the penalties of using HTML socket streams (long-polling and WebSockets) versus TCP streams. HTML socket streams can have up to 5x protocol overhead, up to 3x more payload delivery delay and up to 3x less throughput. In case of small data payload (a few hundreds of bytes), the performance difference between the two becomes significantly noticeable. Also TCP streams tend to behave better over 3G than HTML socket streams. One reason for the poorer performance is due to the fact that the browser uses buffering for HTML socket streams (both long-polling and WebSockets), thus introducing delays, while TCP sockets are able to send the data directly. But on the good side for WebSockets, the paper mentions that they behave better than long-polling when it comes to sending small-sized chunks, making it a better choice for chat, VoIP, online games, etc, just like TCP streams. A big plus for HTML sockets over TCP streams is the fact that the data sent through them can pass over firewalls and most proxies, as HTTP and HTTPS are commonly allowed ports.

Another example where WebSockets are better than AJAX is the case of sending small data per frame. While WebSockets exchange 2B of data per frame, continuous polling with AJAX exchange up to 8 KB of HTTP header [2].

III. EVALUATING EXISTING WEBSOCKET IMPLEMENTATIONS

In order to get a better idea about how different WebSockets implementations behave and how the architecture of the testing application should be designed, a series of tests were performed on two different Node.js [8] WebSocket implementations.


Fig. 1. Test 1 - Constant Flow of Incoming Connections

The chosen implementations are **ws** [9] and **WebSocket-Node** [10], both part of Node.js. The tests were performed on an Intel Core i7-4770 with 16GB of RAM running on Ubuntu 13.10, 64-bit. As a measuring tool we have used the *Look* performance profiler for Node.js [11], which offered information about the process's CPU time, RSS (Resident Set Size), V8 heap total, V8 heap used, and OS memory and load average. We have displayed only CPU time and RSS because we considered them to be the most relevant, but all tests can be found at [12].

We have tested 4 different scenarios:

1. Connected 4 clients simultaneously every second until 3400 connections were reached. All connections are kept alive during the entire run and transferred data is negligible (a single number). Figure 1 contains the comparison between the two WebSockets implementations. We can observe that **ws** has a better CPU usage when handling fewer connections

Fig. 2. Test 2 - Bursts of Incoming Connections

Fig. 3. Test 3 - Constantly Increasing Volume of Data Exchanged Between a Fixed Number of Clients

than **WebSocket-Node**, but towards the end they behave almost similarly. The memory consumption of the two is behaves similarly.

2. Connected 600 clients simultaneously once every 120 seconds. All connections are kept alive during the entire run and transferred data is negligible (a single number). Figure 2 contains the comparison between the two websockets implementations. Following the conclusions of the first example, memory consumption is similar and **ws** is faster than **WebSocket-Node**, but under extremely heavy weight, the system freezes faster for **ws**. We were able to go only up to 3600 connections with **ws**, while **WebSocket-Node** peaked at 4200 connections. .

3. Simultaneously connected 100 clients right away, which exchanged data with sizes ranging from 2KB up to 1MB and increments of 2KB/s per client. All connections were kept

alive during the entire run, and no extra clients were added. Clients simply send more data with every second. Figure 3 contains the comparison between the two WebSockets implementations. In this example we can clearly see that **WebSocket-Node** behaves better when sending large packets. At almost 1MB / packet, **WebSocket-Node** processes everything in under 2s, but **ws** requires up to 6s. Also we had to stop **ws** before 1MB / packet because the system almost froze, while the other implementation allowed us to let it reach its limit. .

4. Simultaneously connected 400 clients, then closed their connections after 10 seconds. Once every 90 seconds, 200 clients would be added to the prior number of clients, and the first steps would be repeated. The process continued until 2800 connections were reached. Transferred data was negligible (a single number). Figure 4 contains the comparison between the two websockets implementations. In this example
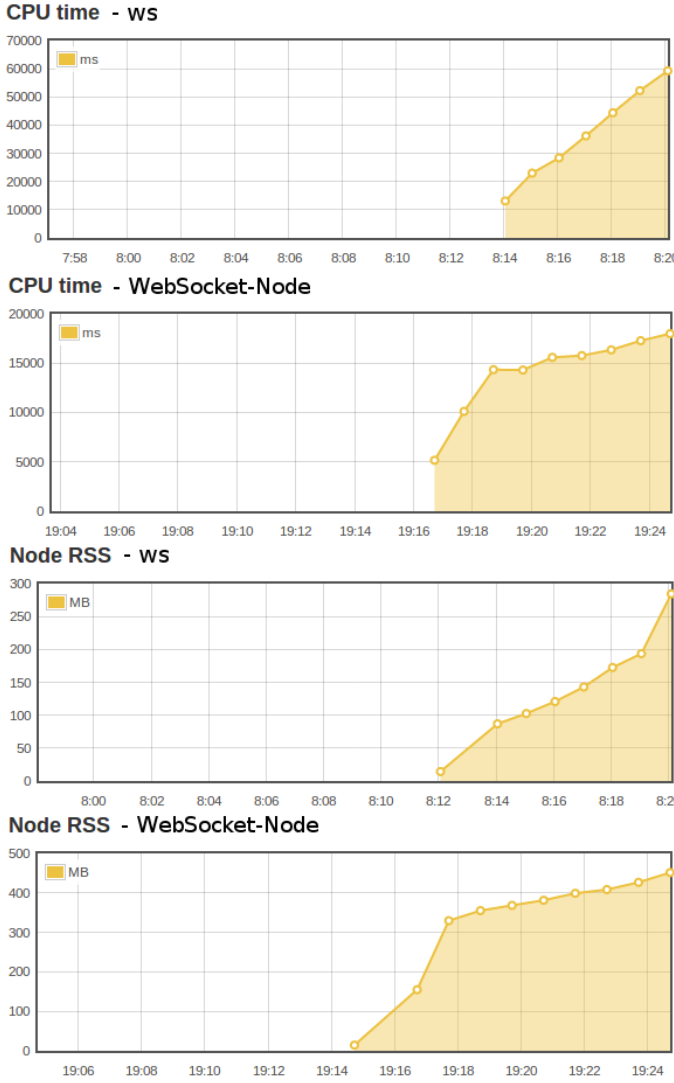
Fig. 4. Test 4 - Linearly Increasing Number of Simultaneous Connections

the used memory is similar, but the **ws** implementation has bigger peaks for CPU usage than **WebSocket-Node** when it comes to the last few iterations, where large numbers of connections must be managed (more than 2200 connections). For small numbers of connections, both implementations behave similar. .

We arrived at the conclusion that **ws** is faster than **WebSocket-Node** when it comes to fewer connections and small data packets, but with larger-sized packets and sizeable amounts of connections, **WebSocket-Node** seems to behave better.

IV. ARCHITECTURE

This section proposes two different architectures for an application which is capable of testing the performance of WebSockets and AJAX implementations of bi-directional communication. The first proposal involves a generic profiling layer over existing WebSockets implementations, coupled with
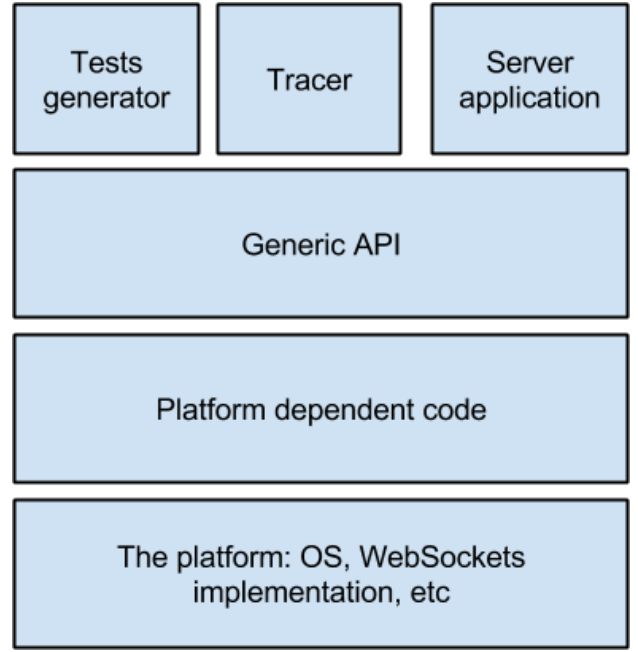

Fig. 5. Testing Framework Based on a Generic API

the necessary modules required in order to extract the useful data. The alternate design involves writing a platform-dependent process analysis tool which is capable of inspecting WebSocket proxies, together with the necessary client implementation needed in order to perform the eventual stress tests.

A. The Generic API Approach

Figure 5 contains the architecture for the first proposal, structured into layers. In order to test the performance automatically, the application should be able to generate tests based on a given input and automatically measure how these tests behave on a given platform. The design should promote easy porting to other devices and Operating Systems. It should have a layer which is dependent of platform and implementation. This layer is the base of a generic API that is used by the tests generator component, the server application which communicates directly with the tests generator, and the tracer used to gather information about the WebSockets behavior inside the server. The API and the previously mentioned components are going to be ported to other platforms with minimal changes. The layer that provides the API is going to be implemented by every platform in a different way, depending on the Operating System, specific WebSockets API requirements, etc. .

B. The Platform-Dependent Alternative

Figure 6 shows the alternate architecture. Just as before, we have a server application and tests generator component that communicate directly. The tests generator creates stress tests based on a user defined input, similar to Unix commands. Similarly to the first architecture, the tracer
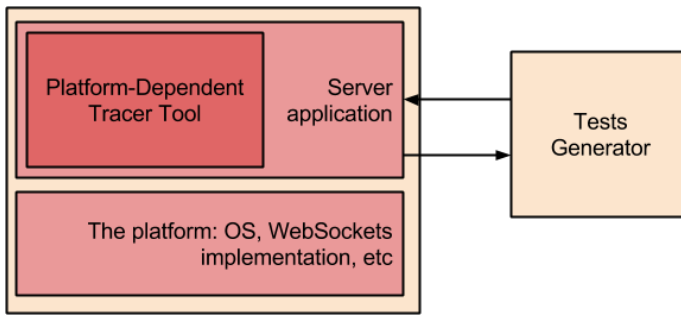
Fig. 6. Testing Framework Based on a Process Monitoring Application

gathers information about the WebSockets behavior while the target server is up and running. However, while in the first case the tracer is a separate component that can gather information independently from the server due to the generic API, in the current architecture the tracer is a library used by the server directly, similar to the Look [11] profiler. As a plus, this offers the possibility to use the tracer in other server applications that are not created only for testing WebSockets implementations in general, but also test the behavior of specific servers that might be used in production and are based on WebSockets. But this comes at the cost of portability. For every new programming language, OS or device, the platform dependent tool and server application need to be completely rewritten. The tests generator does not suffer from this problem because it acts as a client for the server and it will always make use of the same communication protocol. .

V. CONCLUSIONS AND FUTURE WORK

Starting with research on WebSockets and their potential in the web application industry, we continued with a performance evaluation of a couple of existing WebSockets implementations. Having determined the WebSockets performance caps, we were able to propose two possibilities of designing a framework capable of profiling existing web servers and evaluate the benefits of switching to the new connection-oriented WebSockets protocol. One architecture promotes easy porting to new platforms and WebSockets implementations, with the downside that it will initially take more time to implement and can be used only for generic testing. The other is much more straight-forward, as its main purpose is to generate testing results as fast as possible and maybe to be used by applications created for production, but it comes at the cost of harder porting to other platforms.

As future work, we can extend the comparative analysis to other WebSockets implementations, such as **socket.io** [13], **SockJS** [14] or **Engine.IO** [15]. These comparisons might give us better insights about which testing architecture we should choose and how we should implement it, or maybe create a better architecture.

We feel that the performance comparison framework is likely to become very popular if properly written. Effectively seeing the benefits of switching to WebSockets is very likely to accelerate their rate of acceptance, thus making the web an even better, faster and more efficient place.

REFERENCES

[1] Anthony T. Holdener, III. *Ajax: The Definitive Guide*. O'Reilly, first edition, 2008.
[2] Ian Hickson. '[hybi] web socket protocol in "last call" ', 2009. http://www.ietf.org/mail-archive/web/hybi/current/msg00784.html.
[3] A. Melnikov I. Fette. The websocket protocol, 2011. http://tools.ietf.org/html/rfc6455.
[4] Autobahn, testsuite. http://autobahn.ws/testsuite/.
[5] D.G. Puranik, D.C. Feiock, and J.H. Hill. Real-time monitoring using ajax and websockets. In *Engineering of Computer Based Systems (ECBS), 2013 20th IEEE International Conference and Workshops on the*, pages 110–118, 2013.
[6] J. Hill, H. Sutherland, P. Stodinger, T. Silveria, D.C. Schmidt, J. Slaby, and N. Visnevski. Oasis: A service-oriented architecture for dynamic instrumentation of enterprise distributed real-time and embedded systems. In *Object/Component/Service-Oriented Real-Time Distributed Computing (ISORC), 2010 13th IEEE International Symposium on*, pages 10–17, 2010.
[7] S. Agarwal. Real-time web application roadblock: Performance penalty of html sockets. In *Communications (ICC), 2012 IEEE International Conference on*, pages 1225–1229, 2012.
[8] Node.js, 2009. http://nodejs.org/.
[9] Einar Otto Stangvik. ws: a node.js websocket library, 2011. http://github.com/einaros/ws.
[10] Websocket client and server implementation for node, 2011. http://github.com/Worlize/WebSocket-Node.
[11] Vadim Baryshev. Look, 2012. http://github.com/baryshev/look.
[12] Liviu Chircu Mihail Costea. Websockets implementation tests, 2013. https://github.com/cmihail/WSvsHTTP/tree/master/tests.
[13] socket.io, 2010. https://github.com/learnboost/socket.io.
[14] Marek, Serge S. Kova, Matthew Sackman, Simon MacMullen, and Michael Bridgen. Sockjs - websocket emulation, 2011. https://github.com/sockjs.
[15] Guillermo Rauch. Engine.io, 2011. https://github.com/LearnBoost/engine.io.